A Run-Time Reconfigurable Neural Inference Engine with Deterministic Sub-Microsecond Latency on a Low-Cost FPGA SoC

Ahmet Batuhan Özkaner[1*], Mustafa Alper Ilık[1] and Baykal Sarıoğlu[1]

[1*]Department of Electrical and Electronics Engineering, Istanbul Bilgi University, Istanbul, Türkiye.

*Corresponding author(s). E-mail(s): ahmet.ozkaner@bilgiedu.net;
Contributing authors: alper.ilik@bilgiedu.net; baykal.sarioglu@bilgi.edu.tr;

Abstract

Running neural inference on an embedded processor is flexible but leaves latency exposed to scheduling and the software stack; moving it into programmable logic makes latency predictable but normally ties the bitstream to one network, so every retrained model costs a synthesis run. This paper presents an inference engine that avoids both. The network shape is a run-time parameter held in configuration registers, so networks of different topologies are deployed into an unchanged bitstream over AXI-Lite with no synthesis step. On a low-cost Xilinx Zynq-7000 the engine uses 87.3 % of the DSP slices and 27.0 % of the LUTs at an estimated 2.70 W, closes timing at 125 MHz, and completes an inference in 768 ns. An on-chip cycle counter shows this latency to be invariant: over 100 000 inferences the standard deviation is zero. The same network run bare-metal on the SoC's Cortex-A9 requires 52.4 µs, 68 times longer. The datapath was verified bit-exactly against an independent reference model over 400 000 inferences across four topologies, a process that first exposed three pipeline-alignment defects invisible to synthesis. Against the same four networks generated by hls4ml for the same device, clock and arithmetic, over 20 design points, every one that fits placed and routed, no design of the two larger networks meets 125 MHz, and the fastest is 1.3 to 2.0 times slower than the engine even at its own maximum frequency. The engine is evaluated end to end over Ethernet on a latency-critical financial decision task.

Keywords: Embedded systems, FPGA, Run-time reconfiguration, Neural network inference, Deterministic latency

1 Introduction

Embedded systems increasingly need to evaluate a small neural network on every input event, under a latency bound that must hold for the worst case rather than on average. The usual options each give up something. Running the network on the embedded processor keeps the model easy to change but leaves latency exposed to scheduling, interrupt handling, cache behavior and the depth of the software stack, so the distribution acquires a heavy tail even when the mean is acceptable. Moving the network into programmable logic removes those sources of variation, but the standard route through high-level synthesis binds the bitstream to one network: a retrained model, or a change of topology, costs a synthesis and place-and-route run measured in hours.

This paper presents an inference engine that keeps the determinism of the fabric without inheriting that rigidity. The engine runs *inside* the programmable logic of a low-cost Zynq-7000 SoC, is fed over Ethernet through the processing system and an AXI-DMA /

AXI-Stream datapath, and treats the network shape as a run-time parameter: layer widths live in configuration registers, so networks of different topologies are loaded into an unchanged bitstream over AXI-Lite with no synthesis step between them. The input feature set, the fixed-point format and the number of decision classes are configurable on the same terms, which makes the design a reusable embedded inference engine rather than one fixed accelerator.

We evaluate it on a task that makes the latency requirement concrete. In high-frequency trading the time between the arrival of market data and the emission of an order is as important as the correctness of the order itself, and the competition is economically consequential rather than merely technical: Budish et al. [1] model the resulting latency race as an equilibrium feature of continuous limit order books, and Aquilina et al. [2] quantify its cost from exchange message data, finding races decided on the scale of microseconds. The workload used throughout is a three-class directional decision over Binance BTCUSDT trade data, specified in Section 6.3. Nothing in the architecture is specific to it; the task supplies a realistic input rate, a realistic network size and a worst-case requirement that a mean latency figure would not capture.

1.1 Main Contributions

The main contributions of this work are:

- A hardware MLP inference core for a low-cost Zynq-7000 SoC built on 192 DSP-based multipliers and a fabric adder tree in a 13-stage pipeline, whose latency follows an analytic model ($\lceil n_\ell/3 \rceil + 14$ cycles per layer) whose predicted total agrees exactly with the measured input-to-done interval and with an on-chip cycle counter read back on every inference, and whose per-layer predictions agree with the oscilloscope captures to within one clock cycle.
- Run-time reconfigurability of the network *shape*, not only of its weights: layer widths reside in configuration registers, so networks of different topologies are deployed into an unchanged bitstream over AXI-Lite without any synthesis step.
- A three-level verification methodology (an independent bit-accurate reference model, directed RTL simulation, and on-hardware comparison) which uncovered three pipeline-alignment defects invisible to synthesis and to end-to-end testing, and which after correction establishes bit-exact agreement over 400,000 inferences and 1,200,000 raw output values across four topologies.
- An Ethernet-fed hardware/software co-design with a per-stage latency breakdown measured over 10^6 packets, isolating how much of the end-to-end path is attributable to the deterministic hardware core and how much to the host and network software around it.
- Quantitative evidence for the determinism claim rather than illustrative captures: an on-chip cycle counter read back with every inference shows a single-valued latency distribution over 100 000 inferences, and a like-for-like software baseline on the SoC's own Cortex-A9, running identical fixed-point arithmetic, establishes the cost of the alternative.
- A measured comparison against a compiler-generated implementation: the same four networks produced by hls4ml for the same part, clock and bit-exact arithmetic, swept over five reuse factors and taken through place and route, which locates the engine against the compiler-based design space with post-route timing rather than synthesis estimates.

2 Related Work

2.1 FPGA Acceleration in Financial Systems

Reconfigurable logic has been applied to electronic trading since the point at which software network stacks ceased to be competitive at microsecond timescales. Leber et al. [3] moved market-data feed handling and order generation into an FPGA and reported wire-to-wire latencies roughly an order of magnitude below a tuned software path, establishing the pattern that the determinism of a fixed datapath matters as much as its mean latency. Lockwood and Monga [4] pursued the same argument at data-center scale, placing the entire service in programmable logic to remove operating system and kernel-stack variability. More recent work has raised the level of abstraction rather than the level of ambition: Puranik et al. [5] synthesised a trading-system order-processing back end through a high-level synthesis flow, reporting roughly 500 ns per order against approximately 1 µs for the equivalent software, and Brown et al. [6] applied dataflow transformations to a credit-default-swap pricing engine on a data-center accelerator card. LightTrader [7] is the closest system in intent, since it places a learned model inside the trading loop: an FPGA handles the network and order path, while the

inference itself runs on dedicated deep-learning accelerator chips, and a scheduler trades model size against the time remaining before a deadline.

What the fabric-resident designs share is that the accelerated function is fixed at elaboration time. A feed handler, an order book, a pricing kernel or a rule-based strategy is compiled into the fabric; changing the strategy means changing the bitstream. This is appropriate when the algorithm is stable, but it does not address the case in which the decision logic is a *learned* model that is periodically retrained.

2.2 Neural Network Inference on FPGAs

The FPGA deep-learning literature has largely converged on compiler-style frameworks that translate a trained network into a bitstream. Zhang et al. [8] formalised the design-space exploration underlying such accelerators; Qiu et al. [9] showed that the resulting designs fit within the resource envelope of an embedded Zynq part, and Angel-Eye [10] added a quantization flow and a compiler on top of that line of work; FINN [11] and its successor FINN-R [12] generate dataflow architectures for binarized and quantized networks, reaching classification latencies of a few hundred nanoseconds; and hls4ml [13] targets the same sub-microsecond regime for the trigger systems of particle-physics experiments, which resemble the trading problem in that a decision must be produced within a hard latency budget. hls4ml has since matured into a general compiler platform with several model front ends and hardware back ends [14], and da4ml [15] replaces its fully unrolled constant matrix-vector products with distributed arithmetic, cutting on-chip resources by up to a third while also reducing latency. Both gain their efficiency from the same source: the weights are constants of the circuit. The same low-cost XC7Z020 used here has also been targeted through Vitis HLS for latency-bound signal processing, meeting a 70 µs worst-case detection budget for LTE-A cognitive radio [16]. Two recent contributions to this journal show the range of the design style on Zynq-class parts. Gdaim et al. [17] drive an OpenCL high-level synthesis flow on a Zynq-7000 with predictive models of loop latency, throughput and BRAM feasibility that are evaluated before the hardware is deployed, and Rangsikunpum et al. [18] implement a binarized traffic-sign classifier with an adversarial detector on Zynq SoCs including the ZedBoard used here, at over 9,000 frames per second within about 2.8 W. Both are evaluated, as is usual for vision workloads, in frames per second rather than in single-inference latency. SparkNoC [19] and comparable edge accelerators optimize energy efficiency for convolutional workloads, while Tsoutsouras et al. [20] show the same exploration methodology applied to a non-neural classifier on a Zynq SoC; Wang et al. [21] survey the field more broadly. On embedded SoC targets the movement of data between the processing system and the fabric is frequently the limiting factor rather than the arithmetic itself, as quantified for a CNN accelerator by Rios-Navarro et al. [22], an observation that the per-stage breakdown of Section 6 reproduces for the present design.

The arithmetic these designs rely on is well understood. Fixed-point and integer-only inference has been shown to preserve accuracy for a wide range of networks [23, 24], and the trade-offs are surveyed comprehensively by Gholami et al. [25]; for multi-layer perceptrons specifically, the propagation of quantization noise through the layers can be predicted analytically [26]. The $Q8.8$ format adopted here sits at the conservative end of that spectrum.

The common property of this second group is that the network topology is baked into the generated hardware. FINN, FINN-R, Angel-Eye and hls4ml all require a new synthesis and implementation run when the network changes, which for a mid-size design is a multi-hour turnaround. They are also, without exception, evaluated as compute kernels: the input arrives from host memory or an on-chip stream, and the network interface is outside the scope of the reported latency.

2.3 Run-Time Programmable Overlays

A separate line of work removes the re-synthesis requirement by placing a fixed, programmable engine on the fabric and expressing the network as data rather than as logic. VTA [27] defines an instruction set for tensor operations and a compiler stack above it, so that a change of network is a change of instruction stream. OPU [28] and its lightweight successor [29] make the same argument specifically for convolutional networks on FPGAs, and state the property explicitly: the network is compiled to microcode that the overlay executes without reconfiguration. Commercial toolchains adopt the same model: the deep-learning processing unit shipped by AMD for Zynq and Versal parts [30] is a fixed overlay driven by compiled instructions, so that a new network is deployed by

recompiling its instruction stream rather than by re-implementing the device. Khaki and Choi [31] use that flow to deploy pruned and quantized lightweight CNNs and report latencies of 6.28 to 19.95 ms per frame, which is the operating point such overlays are built for. The idea has since been carried to transformers: the runtime-adaptive accelerator of Kabir et al. [32] executes different transformer models through software-set parameters without re-synthesis. Ratto et al. [33] derive run-time reconfigurable accelerators from an ONNX description, but use the reconfigurability to trade execution time for power within one network rather than to change the network. Closest to the present work in its latency target, SLAC's Neural Network Library [34] keeps HLS-generated layers in the style of hls4ml, against which it is benchmarked, but loads weights and biases at run time without re-synthesis; the published description presents this as a reload of the weights rather than as a change of the layer dimensions.

Dynamic partial reconfiguration offers a third mechanism, in which the fabric itself is rewritten at run time. Boudjadar et al. [35] apply it to CNN acceleration on the same XC7Z020 used here, reporting 0.68 s per image. Partial reconfiguration avoids a full synthesis run at deployment, but each module still has to be synthesized in advance and the reconfiguration itself takes the device region out of service, which rules it out for changes made while inference is in progress.

The engine described in this paper belongs to this family, and we make no claim to have introduced run-time programmability. The distinction lies in the design point rather than in the mechanism. Published overlays are throughput architectures: they are dimensioned for convolutional workloads, evaluated in frames or images per second, and their reported latencies are per batch or per frame rather than per inference. Among published overlays that change the network shape without re-synthesis we are aware of none that reports single-inference latency below one microsecond, and the reason is structural rather than incidental: overlay evaluations are reported per frame or per batch, and the tiled dataflow that gives an overlay its generality implies that a network of this size occupies the engine for many passes. We state the claim in this restricted form because it rests on the absence of a reported figure rather than on a measurement of our own. The present design accepts a much narrower scope, fully connected layers within fixed bounds, in exchange for a latency figure that is a property of a single inference.

2.4 The Latency Extreme

At the opposite end of the trade-off, LogicNets [36] abandons arithmetic units altogether and compiles quantized neurons directly into lookup-table logic, reaching inference latencies of the order of nanoseconds for networks comparable in size to the one considered here. Subsequent work in the same direction has extended the approach to higher-order functions, by learning piecewise polynomials inside each lookup table [37], and to deeper topologies, by hiding a multi-layer sub-network inside a single synthesizable Boolean function [38]. More recent work widens the inputs each table can take [39], assembles larger networks from LUT-mapped sub-networks [40], and applies the same mapping to gradient-boosted decision trees [41].

These results are one to two orders of magnitude faster than the 768 ns reported below, and we state that plainly rather than selecting a comparison that flatters the present work. The relevant question is what the difference buys. Compiling a network into LUT logic makes its weights part of the circuit: retraining requires a new synthesis, place and route, and bitstream. For a model that is retrained on a schedule, that turnaround is the binding constraint rather than the inference latency. The design space therefore has at least three distinguishable points, and the present work occupies the middle one: LUT-mapped networks give nanoseconds with per-model synthesis, overlays give re-programmability with per-frame throughput, and the engine described here gives sub-microsecond single-inference latency with re-programmability, on a low-cost part.

2.5 Learned Models for Order-Book Prediction

On the modelling side, deep networks applied to limit-order-book data are an established line of work. Zhang et al. [42] combine convolutional and recurrent layers for mid-price movement prediction, and the FI-2010 benchmark of Ntakaris et al. [43] provides a common evaluation basis. This literature is concerned with predictive quality and is typically evaluated offline; the deployment latency of the resulting model is rarely reported, and the models are generally too large for a microsecond decision budget without substantial compression.

2.6 Positioning of This Work

Taken together, the preceding sections describe a trade-off between how fast a single inference completes and how expensive it is to change the model. LUT-mapped networks sit at one end, compiler-generated dataflow accelerators in the middle, and overlays at the other end. What is missing from that spectrum, and what this paper contributes, is a point that combines sub-microsecond single-inference latency with re-programmability, on a part that costs a few hundred dollars, together with evidence that the deployed arithmetic is correct.

We therefore do not claim run-time reconfigurability as such, nor the lowest latency in the literature. The claim is the conjunction of four properties, none of which is individually novel but whose combination we have not found reported: (i) single-inference latency below one microsecond, given as an analytic function of the network shape and confirmed on hardware; (ii) deployment of networks of different topologies into an unchanged bitstream; (iii) an Ethernet to decision path measured end to end, with the contribution of each stage separated; and (iv) bit-exact agreement between the deployed arithmetic and an independent reference model over 4×10^5 inferences. The fourth is uncommon in this literature and, as Section 6.2 reports, was not a formality in our case.

Table 1 positions the work against representative prior designs. We note that the comparison of absolute latency figures across these rows is of limited value, since the reported quantities are not measured between the same two points: FINN and hls4ml report a compute-kernel latency, whereas the figure reported here is the end-to-end path from an Ethernet frame to a decision. The core latency of this design is given separately in Section 6 to make the comparison meaningful.

For the compiler-based part of the design space we do not rely on these figures alone. Section 6.14 reports a like-for-like implementation of the same four networks generated by hls4ml on the same device, at the same clock and with bit-identical arithmetic, so that this end of the spectrum is represented by a measured point rather than an argued one.

3 Proposed Framework

3.1 System Overview

The proposed system follows a strict hardware/software partition on a single Zynq SoC, illustrated in Fig. 1. Market data arrives over Ethernet and is received by the PS, which prepares the feature vector and issues it to the PL through an AXI-DMA engine. The custom `mlp_axi_wrapper` IP receives the stream, runs the inference core, and returns the decision and status through AXI-Lite. The PS therefore performs only I/O and control, while the entire inference workload is executed in the PL fabric. The end-to-end conceptual flow is:

$$\text{Ethernet} \rightarrow \text{PS} \rightarrow \text{AXI-DMA/Stream} \rightarrow \text{MLP core} \rightarrow \text{BUY/HOLD/SELL}.$$

3.2 Ethernet-Enabled Data Path

Feature packets are delivered over Ethernet to the PS, parsed, and forwarded to the PL. Each 32-bit AXI-Stream word carries two 16-bit $Q8.8$ values, and the 40 features, zero-padded to the 64-wide datapath, are transferred as 32 stream words on which the engine auto-starts. Every stage of this path is instrumented with PS-side timestamps and quantified over a 1,004,229-packet run in Section 6 (Table 12).

3.3 Customization Flow

The framework is reusable across models and strategies. A user trains an MLP on their own dataset and feature set; the trained parameters are quantized on the host and written, together with the layer widths, into the on-chip weight and bias memories over AXI-Lite, while the bitstream, and with it the datapath, AXI interfaces, and control logic, is unchanged. Table 2 summarizes the configurable components.

4 FPGA-Embedded Neural Inference Architecture

The embedded network is an MLP with topology $40 \rightarrow 64 \rightarrow 32 \rightarrow 16 \rightarrow 3$: a 40-dimensional input feature vector, three hidden layers with ReLU activation, and a 3-class output (BUY/HOLD/SELL); the argmax over the three raw scores is evaluated by the PS after readback. Inference is performed entirely in hardware; no layer is offloaded to the PS.

Table 1 Comparison with representative FPGA-based financial-acceleration and neural-network inference studies. "Re-synth." indicates whether changing the deployed model requires a new bitstream. The reported metrics are not measured between the same two points and are given for orientation rather than as a ranking; see the discussion in the text.

Study	Platform	Function	Precision	Re-synth.	Net. I/O	Reported metric
Leber et al. [3]	Virtex-5	Feed handler, order gen.	n/a	Yes	Yes	2.6 μs half round trip
Puranik et al. [5]	Zynq	Order processing, HLS	n/a	Yes	No	~500 ns per order
Brown et al. [6]	Alveo U280	CDS pricing	Float	Yes	No	27,675 options/s
FINN [11]	ZC706	Binarised MLP/CNN	1 bit	Yes	No	0.31 μs, 12.3 M cls./s
hls4ml [13]	Virtex-7	MLP, trigger	Fixed	Yes	No	~100 ns, jet tagging
Angel-Eye [10]	Zynq XC7Z045	CNN compiler	8/16 bit	Yes	No	137 GOP/s, VGG16
SparkNoC [19]	Arria 10 GX1150	Lightweight CNN	Fixed	Yes	No	337.2 GOP/s
VTA [27]	Zynq, Ultra96	Tensor overlay	8 bit	No	No	Per-frame throughput
OPU [28]	Zynq XC7Z100	CNN overlay	8 bit	No	No	Per-frame throughput
LogicNets [36]	Virtex UltraScale+	LUT-mapped MLP	2 to 6 bit	Yes	No	~10 ns
This work	Zynq XC7Z020	MLP, run-time shape	16 bit	No	Yes	768 ns core, 68.6 μs end to end

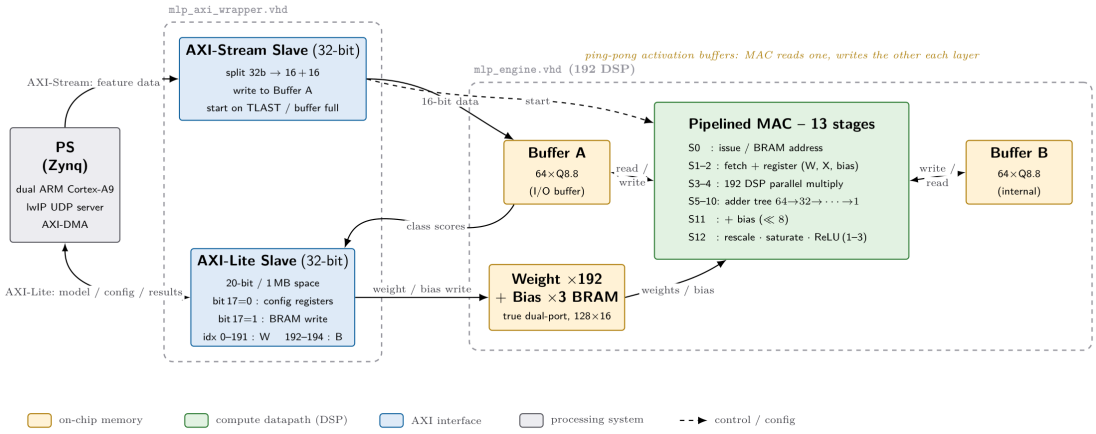


Fig. 1 System architecture of the proposed framework on the Zynq-7000 SoC. Market data arrives over Ethernet at the PS, which forwards the prepared feature vector to the custom `mlp_axi_wrapper` core through AXI-DMA / AXI-Stream; the model parameters are loaded and the final decision is read back over AXI-Lite. The entire inference is executed by the 13-stage pipelined inference engine in the PL.

4.1 Fixed-Point Quantization

The floating-point model is converted to a $Q8.8$ fixed-point representation (8 integer, 8 fractional bits), which lowers resource usage and enables single-cycle DSP multiplications while preserving accuracy. Weights and biases are stored in BRAM and addressed through the AXI-Lite configuration interface. Table 3 reports the effect of quantization on the reference model. The fidelity of the *hardware* implementation of this arithmetic is a separate question and is established independently in Section 6.2.

4.2 Parallel-Multiplier Inference Core

The inference core (`mlp_engine`) is built around 192 DSP slices used as parallel multipliers. The accumulation is not performed inside the DSP blocks but by a fabric adder tree, which is why the logic utilization is dominated by the reduction network rather than by the multiplications. The core is organized as a 13-stage pipeline (Fig. 2): one address-issue stage (S0) followed by twelve datapath stages. Stages 1 and 2 read weights/inputs from BRAM and register them; stages 3 and 4 perform the 192 parallel multiplications; stages 5 to 10 reduce the products through a balanced $64 \rightarrow 1$ adder tree; stage 11 adds the bias; and stage 12 rescales

Table 2 Customization options supported by the proposed framework

Component	Customizable Parameters
Input features	Feature vector contents and size (default 40)
MLP model	Retrainable weights/biases; hidden-layer sizes
Fixed-point format	Word/fraction split (default $Q8.8$)
Output decision	Class set, e.g. BUY / HOLD / SELL
Deployment flow	Host-side quantization and AXI-Lite upload, no synthesis

Table 3 Effect of $Q8.8$ quantization on the reference model

Metric	Value
Float accuracy	61.74%
$Q8.8$ accuracy	61.65%
$Q8.8$ vs. float agreement	99.93% (macro F1 0.990)
Class change ratio (float vs. $Q8.8$)	0.074%

the result and applies the ReLU activation before write-back. Two ping-pong buffers (Buffer A / Buffer B) hold the inter-layer activations so that successive layers reuse the same multiplier array.

5 Implementation on Zynq-7000 SoC

5.1 Hardware/Software Partitioning

Table 4 lists where each function is implemented. The PL implements the entire inference datapath and the measurement signals; the PS is responsible for the network interface, for handing the received feature vector to the DMA engine, and for high-level control.

Two aspects of this partition should be stated plainly, because they bound what the reported latency covers. First, feature construction and the conversion to $Q8.8$ are performed *upstream of the board*: the packet arriving at the PS already contains a quantized, zero-padded vector, and the PS copies it into a DMA-aligned buffer without transforming it. Moving feature construction onto the PS would add work to the measured path and is not claimed here. Second, the argmax over the three output scores is evaluated in PS software after the scores have been read back over AXI-Lite; it is a three-way comparison and is included in the read-out stage of the latency breakdown (Table 12) rather than in the core latency. Moving it into the PL would remove two of the three AXI-Lite reads and is a straightforward extension.

Table 4 Hardware/software partitioning of the proposed system

Component	Implementation Location
Ethernet packet reception	PS (software)
Feature vector handoff	PS (software, `memcpy`)
Feature vector transfer	AXI-DMA (PS→PL)
MLP inference	PL (hardware)
Argmax decision	PS (software, after read-back)
Status/result registers	AXI-Lite (PS↔PL)
Debug signals	PL output pins

5.2 AXI-Stream and AXI-DMA Interface

The PS prepares the feature vector and transfers it to the PL with AXI-DMA; the `mlp_axi_wrapper` consumes it as an AXI-Stream slave, splitting each 32-bit word into two $Q8.8$ operands and starting the engine once a full vector is received. Control and status (start, busy, done, result) use AXI-Lite. This keeps the inference pipeline continuously fed without per-operand PS intervention.

5.3 Debug and Measurement Infrastructure

To measure hardware timing directly, the core exports the dedicated debug signals of Table 5. `dbg_data_in_pulse` marks the arrival of a complete feature vector, `dbg_layer1_done` to `dbg_layer4_done` pulse as each layer finishes, and `dbg_mlp_busy` brackets the whole inference. These pins are probed with an oscilloscope, so the per-layer latency is observed on real silicon rather than estimated (Section 6).

6 Experimental Evaluation

6.1 Experimental Setup

The framework is implemented on a Xilinx Zynq-7000 SoC. A host control application (Fig. 3) connects to the C++ host engine, uploads the quantized model, and streams either real CSV market data or randomized test packets while reporting the last prediction, last latency,

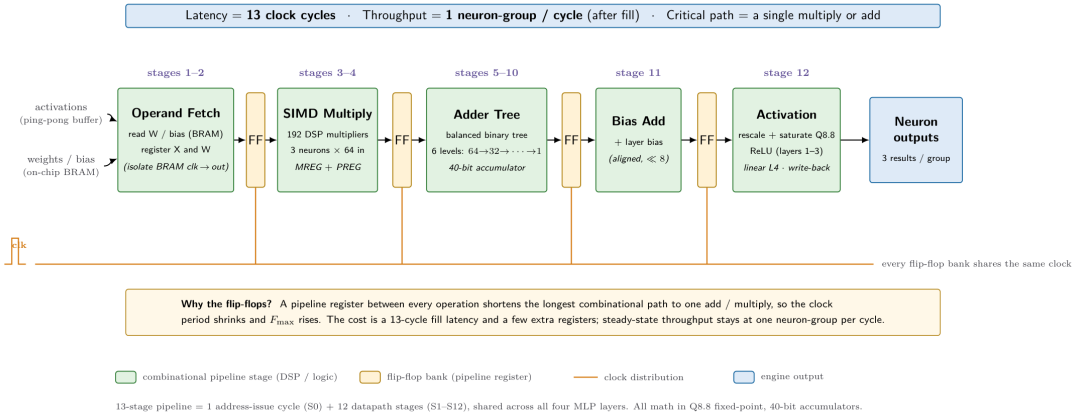


Fig. 2 Pipeline datapath of the hardware MLP core (`mlp_engine`). After the address-issue stage (S0), stages 1 and 2 fetch the weights, inputs, and bias from BRAM; stages 3 and 4 perform the 192 parallel $Q8.8$ multiplications on the DSP slices; stages 5 to 10 form the balanced $64 \rightarrow 1$ adder tree; stage 11 adds the bias; and stage 12 rescales and applies the ReLU before the result is written back to the ping-pong activation buffers. The `mlp_axi_wrapper` splits each 32-bit AXI-Stream word into two 16-bit $Q8.8$ values and triggers the engine; weights/biases reside in BRAM and are loaded via AXI-Lite (address bit 17 selects configuration vs. memory).

Table 5 Debug signals exported for hardware latency measurement. All signal names carry the prefix `dbg_`, omitted here for space.

Signal	Purpose
`data_in_pulse`	Full feature vector received
`layer1_done`	Layer 1 ($40 \rightarrow 64$) complete
`layer2_done`	Layer 2 ($64 \rightarrow 32$) complete
`layer3_done`	Layer 3 ($32 \rightarrow 16$) complete
`layer4_done`	Layer 4 ($16 \rightarrow 3$) complete
`mlp_busy`	Inference in progress; decision valid on the falling edge

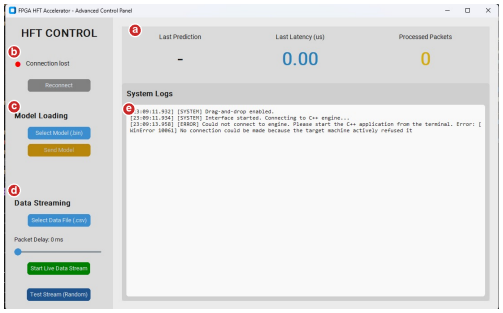


Fig. 3 Host control application: (a) live decision, round-trip latency and packet count, (b) engine connection state, (c) model upload, (d) stream source and inter-packet delay, (e) event log. Deploying a different network is the single action in panel (c); no synthesis step is involved.

and processed-packet count. Table 6 summarizes the setup.

Two kinds of run are reported in what follows and should not be conflated. The first measures the host round trip on the bitstream that is characterized throughout the paper: the host timestamps a request before transmission and after the reply arrives, over 100 000 requests for each of the four networks of Section 6.2, from a native Linux host and, for comparison, from a Windows host. These runs are the basis of Table 11 and of Section 6.6. The second is a longer instrumented run of 1,004,229 packets, in which the firmware timestamps each stage of the on-board path with the Global Timer and returns the timestamps in the reply; it is the basis of Table 12. The two give slightly different end-to-end figures for the reference network (a median of 66.9 µs against 68.5 µs) because they were taken at different times with different host load; the on-board subtotal, which both runs record, agrees to within 0.4 µs in the median. The difference is small relative to the host-side variability characterized in Section 6.6, and no claim in this paper depends on it; where a single end-to-end number is quoted we use the instrumented run, since it is the one whose composition is known.

6.2 Functional Verification of the Inference Core

Reporting latency and resource figures for an accelerator is only meaningful if the accelerator is first shown to compute the intended function. We therefore verify the

Table 6 Experimental setup

Item	Description
FPGA device	Xilinx XC7Z020-1CLG484 (ZedBoard)
Design language	VHDL (PL); C++/Python (PS/host)
Tool version	Xilinx Vivado / Vitis 2025.2
PS software	Bare-metal, lwIP raw API (no OS on the board)
Data path	Ethernet → PS → AXI-DMA → AXI-Stream → PL
Inference core	MLP 40-64-32-16-3, $Q8.8$, 192 DSP multipliers
Dataset	Binance BTCUSDT aggregated trades, 16 631 211 records, 2025-05-01 to 2025-05-18
Task	3-class direction, horizon 50 trades, $\theta = 5 \times 10^{-5}$
Validation split	Final 10 % chronologically, 1 663 113 samples
Host round-trip runs	$4 \times 100\,000$ requests per host (native Linux, Windows)
Instrumented run	1,004,229 packets (PS-timestamped, per stage)
Clock frequency	125 MHz (PL)

core at three levels before reporting any performance result.

Golden reference model.

A bit-accurate software model of the datapath was written independently of the RTL. It reproduces every arithmetic detail of the deployed pipeline: the host $Q8.8$ conversion with saturation, the truncating weight conversion performed by the on-board loader, the 40-bit accumulation, the arithmetic right shift, the 16-bit saturation, the ReLU on all but the output layer, and the argmax rule used by the decision logic. A second, cycle-accurate model of `mlp_engine` was also written; it reproduces the issue FSM, the 13-stage token pipeline, the one-cycle BRAM read latency and the ping-pong write-back, and it is used to reason about pipeline alignment independently of the simulator.

Directed RTL simulation.

Three directed tests exercise mutually independent mechanisms of the engine: a pass-through chain that isolates the multiply, reduce and rescale datapath, a bias-only chain that isolates the bias pipeline, and a multi-neuron test that isolates the write-back indexing. Each test writes its parameters through the same AXI-Lite address map used by the on-board loader, streams a stimulus vector over AXI-Stream, and compares the three read-back outputs against values computed by hand. Input positions that are not overwritten by the network are pre-loaded with distinctive values, so that a missing write-back cannot be mistaken for a correct result.

This step exposed three pipeline-alignment defects in the original design that were invisible to synthesis and to end-to-end testing (Table 7). All three are one-line index corrections; none changes the pipeline depth, the resource usage or the timing closure of the design. After correction, the cycle-accurate model matched the golden reference on 240 randomly generated cases (60 random parameter sets × 4 random input vectors), whereas the original alignment matched on none.

Table 7 Pipeline-alignment defects found by directed simulation

Signal	Stage	Consequence
Bias pipeline	$11 \rightarrow 9$	Each neuron received the bias of a neuron six positions earlier
Write-back token	$12 \rightarrow 11$	Results were written to the wrong activation indices
Input multiplexer	$2 \rightarrow 1$	Each layer read from the wrong ping-pong buffer

On-hardware bit-exact verification.

The corrected design was verified on the board against the golden model over a stimulus set constructed to exercise the saturation, ReLU, sign-transition and zero-padding paths, rather than over recorded market data. The reason is that the features of Section 6.3 are clipped to $[-5, 5]$ and therefore occupy a narrow band near the center of the $Q8.8$ interval: replayed market data never drives the accumulator toward saturation and never exercises the clamping logic, so it cannot distinguish a correct datapath from one whose saturation path is broken. A verification stimulus has to reach the corners that the application itself does not. The firmware was extended to return the three raw 16-bit output values alongside the decision, so that the comparison is made on the arithmetic itself and not only on the argmax.

Table 8 reports the result. Across four networks the hardware reproduced the golden model on 400,000 inferences and on all 1,200,000 raw output values, without a single discrepancy. We note that agreement on the decision alone would be a weak check when the

output distribution is concentrated on one class; the per-output comparison is reported for this reason.

A second weakness is specific to a reconfigurable engine and has to be closed explicitly: if the loader and the reference model both mapped every network onto the same fixed shape, for example by padding or clipping it to the largest layer widths, they would agree perfectly while the shape mechanism itself went untested. The campaign is constructed to exclude this. The golden model evaluates each network at its true, untruncated topology; the loader writes the per-layer width registers of Section 6.13 before streaming the stimulus; and the cycle counter of Section 6.8, logged with every inference, records how many iterations the controller actually issued. The last column of Table 8 shows a distinct, single-valued count for each topology, each one cycle below the value (1) predicts for that shape, the offset being the counting convention explained in Section 6.8. Had the hardware run a padded network, all four counts would have been equal. The four runs used the bitstream whose resource, timing and power figures are reported in Section 6.11; the bitstream inside the boot image was compared byte for byte against the implementation output.

Table 8 On-hardware bit-exact verification against the golden model. All four networks were loaded at run time into the *same* bitstream, without a reset between them. "Outputs" compares the three raw 16-bit outputs of every inference, i.e. 1,200,000 individual comparisons in total, none of which differed. "Cycles" is the on-chip counter value, identical on every inference of a run and equal to $\sum_\ell \lceil n_\ell/3 \rceil + 55$.

Network	Inferences	Decision	Outputs	Cycles
40-64-32-16-3	100,000	100.0000%	100.0000%	95
40-64-64-32-3	100,000	100.0000%	100.0000%	111
40-32-16-8-3	100,000	100.0000%	100.0000%	76
40-10-16-8-3	100,000	100.0000%	100.0000%	69
Total	400,000	100.0000%	100.0000%	

6.3 Dataset and Classification Task

The workload used throughout this paper is a three-way directional classification on Binance BTCUSDT aggregated trades. The data are the public aggregated-trade archive for the instrument, covering 16 631 211 trades between 2025-05-01 00:00 UTC and 2025-05-18 00:44 UTC, approximately seventeen days of continuous tick data. Each aggregated trade record carries a price, a quantity and a taker-side flag.

The 40 features are derived from a trailing window of 40 trades and are dimensionless by construction, which is what makes a $Q8.8$ datapath appropriate. They comprise 20 slots holding successive relative price changes scaled by 1000, 10 trade quantities each divided by the mean quantity of the same window, and 10 taker-side indicators taking the values $+1$ and -1. Only 19 of the 20 price-change slots carry data; the remaining one is a constant zero, retained from an implementation in which the difference series was zero-padded to preserve length. It carries no information, and we report it rather than defend it as a design choice. Every feature is clipped to $[-5, 5]$, so the integer part of $Q8.8$ carries the dynamic range of the input with more than four bits of headroom, and the eight fractional bits determine the resolution.

The label is defined on a forward horizon of 50 trades. Writing P_t for the price of trade t and $r_t = (P_{t+50} - P_t)/P_t$, the class is *buy* for $r_t > \theta$, *sell* for $r_t < -\theta$ and *hold* otherwise, with $\theta = 5 \times 10^{-5}$. The threshold is a design choice that sets the balance of the problem, and it was selected because the two alternatives examined, 10^{-4} and 2×10^{-4}, leave the minority classes at 5 % and below 0.1 % of the validation set respectively, at which point the task degenerates into predicting the majority class.

The split is temporal and not random: the first 90 % of the chronologically ordered samples are used for training and the final 10 %, 1 663 113 samples, for validation. No sample in the validation period is available to the model at training time, so the reported figures are not affected by look-ahead leakage. No embargo was applied at the boundary, so the labels of the last 50 training samples are formed from prices that fall inside the validation period; at 50 samples out of 1 663 113 the effect is not measurable, but we note it rather than leave it to be inferred. The validation period contains 57.16 % hold, 21.50 % buy and 21.35 % sell, which fixes the majority-class baseline at 57.16 %. That figure, and not 33 %, is the reference against which the accuracies below must be read.

6.4 Model Accuracy and Classification Performance

Table 10 reports, for five trained networks, the accuracy in floating point under PyTorch, the accuracy of the same weights through the fixed-point datapath, and the margin over the majority-class baseline established above.

Two properties of the table are worth separating. The first concerns the arithmetic: floating point and $Q8.8$ agree to within 0.22 percentage points across all five networks and to within 0.1 percentage points for the reference network, so reducing the datapath to 16-bit fixed point does not materially change the decisions the network makes. The second concerns the task: the reference network exceeds the majority-class baseline by 4.58 percentage points on 1 663 113 held-out samples, whereas the two smallest networks do not. At this sample size the standard error of a proportion is 0.04 percentage points, so the margins of the three larger networks are two orders of magnitude beyond sampling noise, while network D's margin of 0.06 points is not distinguishable from zero and network E is measurably below the baseline. The classification problem is therefore learnable at this feature set but only above a capacity threshold, and we report the two networks that fail as readily as the three that succeed, since a table containing only the successful configurations would misrepresent how narrow the margin is.

The fixed-point column is computed with a bit-accurate model of the deployed datapath rather than by streaming 1 663 113 samples through the board, which would take hours over UDP. This is legitimate here for a specific reason: Section 6.2 establishes that the hardware and that model agree bit for bit, on every decision and on every raw output value, over 400 000 inferences across four topologies. The fixed-point column is therefore a hardware result in the sense that matters, and the verification section is what licenses reading it that way.

Table 9 gives the per-class behavior of the reference network, which the aggregate accuracy conceals. Recall is roughly balanced across the three classes, between 0.59 and 0.64, so the network is not simply defaulting to the majority class; precision, however, is markedly lower on the two minority classes, around 0.51 to 0.57 against 0.70 for hold. In trading terms the network identifies most directional moves but pays for it with false positives, which is the expected shape for a threshold this small.

Predictive quality is not a contribution of this work, and no claim is made that this network would be profitable after transaction costs. It is reported here so that the workload the hardware executes is fully specified, and so that the fixed-point agreement in Table 10 is anchored to a task whose difficulty the reader can judge.

Table 9 Per-class performance of the reference network ($40 \to 64 \to 32 \to 16 \to 3$) on the 1 663 113-sample validation period

Class	Share	Recall	Precision	F_1
Hold	57.16%	0.619	0.698	0.656
Buy	21.50%	0.595	0.570	0.582
Sell	21.35%	0.637	0.506	0.564
Macro average		0.617	0.591	0.601

Table 10 Floating-point against $Q8.8$ accuracy on the 1 663 113-sample validation period, and the margin over the 57.16 % majority-class baseline. Model A is the reference network used throughout the paper.

Model	Topology	Float (%)	$Q8.8$ (%)	$\lvert\Delta\rvert$ (pp)	vs. base (pp)
A	40-64-32-16-3	61.74	61.65	0.09	+4.58
C	40-64-48-32-3	60.62	60.54	0.08	+3.46
B	40-64-64-32-3	60.58	60.59	0.02	+3.42
D	40-32-16-8-3	57.10	56.88	0.22	−0.06
E	40-10-16-8-3	56.86	57.01	0.15	−0.30

6.5 System-Level Packet-to-Decision Latency

Table 11 reports the measured end-to-end packet-to-decision latency of the full FPGA system from a native Linux host, over 100 000 requests for each of the four networks, all deployed into the same bitstream in one session. Each request is a UDP round trip from the host to the board and back, so it includes the network, the PS software, the AXI-DMA transfer, and the PL core. For the reference network the median is 66.9 µs; the pure hardware core (768 ns) is therefore about 1 % of the end-to-end time, confirming that the remaining cost is dominated by the operating system, the network, and the PS software rather than by the inference itself.

The four columns differ on the board only where the architecture says they should. The median on-board subtotal rises from 9.48 µs to 9.98 µs across the four networks, entirely in the PL stage and in step with the cycle counts, which the firmware reads at the 250 ns granularity of its AXI-Lite polling; reception, parsing, DMA and readout are identical to within 0.01 µs. The end-to-end medians of the three larger networks lie within 0.7 µs of one another. The smallest network, run first in the session, has a median 5 µs lower, and the cause is not on the board. Its run places 38 % of

requests below 60 µs, against 3 % for the other three, and in that fast group the on-board time is the same 9.48 µs as in the slow group; the whole difference, 44.7 against 51.8 µs of off-board time, is spent in the host, its NIC and the wire. We therefore read the lower median as a property of the host state during that run, not of the network, and the variation between columns as a measure of how much the host alone moves the end-to-end figure between runs of a few seconds.

Table 11 Measured end-to-end packet-to-decision latency of the FPGA system from a native Linux host, 100 000 requests per network, all four networks in one unchanged bitstream (all values in μs). The board row is the median of the per-packet on-board subtotal (PS + PL); the cycle row is the core cycle count read back with every inference.

Hidden layers	10-16-8	32-16-8	64-32-16	64-64-32
Minimum	52.4	52.5	53.2	53.6
Median	61.5	66.4	66.9	67.1
Mean	62.7	67.6	68.2	68.2
99th percentile	81.3	79.8	80.9	81.8
99.9th percentile	146.1	116.2	113.6	112.8
Maximum	3260.4	3265.0	3260.9	3244.4
Board median	9.48	9.50	9.75	9.98
Core cycles	69	76	95	111

6.6 How Much of the End-to-End Latency Belongs to the Host

The same UDP round-trip run was repeated from a Windows host with the board, the bitstream and the four networks held fixed, which bounds how much of the end-to-end figure is determined by software outside the board. For the reference network the Linux host completes a median request in 66.9 µs (p99 80.9 µs, p99.9 113.6 µs); the Windows host needs 460.0 µs (p99 656.0 µs, p99.9 1878 µs), and 255 of its 100 000 requests exceed 1 ms, against 30 on Linux. The Windows medians of the four networks lie between 456.1 and 461.0 µs, so on that host the network shape is not visible in the end-to-end figure at all. A factor of roughly seven in the median, and the entire heavy tail, therefore originates on the host side of the link rather than in the accelerator. This is also why the end-to-end figures in Table 11 should not be read as a property of the core. This experiment is not a software baseline for the inference itself, since the network is executed in the fabric in both runs; that comparison is given in Section 6.10.

6.7 Latency Breakdown

Because each layer exposes a done pulse, the inference can be decomposed on real hardware. Probing the debug pins with an oscilloscope yields the per-layer inter-pulse times of Fig. 4 and the total input-to-final-done interval of Fig. 5. Table 13 converts these to clock cycles at 125 MHz. The measured per-layer times are 280, 200, 160, and 118 ns for layers 1 to 4, and the measured end-to-end core interval is 768 ns.

These values are not merely observed; they follow from the architecture. The issue FSM emits one token per cycle while a layer is being issued, and each token carries three neurons, so a layer of n_ℓ outputs occupies $\lceil n_\ell/3 \rceil$ issue cycles. The FSM then waits for the pipeline to empty before advancing, which costs the full pipeline depth. The per-layer occupancy is therefore

$$c_\ell = \left\lceil \frac{n_\ell}{3} \right\rceil + D + 1, \qquad D = 13, \tag{1}$$

where D is the depth of the token pipeline. Analytic latency models of this kind are also used to steer design-space exploration for systolic accelerators, where they have been validated against FPGA measurements to within about one percent [44]; for a single-issue pipeline without data-dependent control the model can be exact. Table 13 compares this prediction with the oscilloscope measurement. The predicted total, $\sum_\ell c_\ell = 96$ cycles or 768 ns, agrees exactly with the measured input-to-done interval. Two of the four individual intervals read slightly below the prediction, layer 1 at 280 ns against a predicted 288 ns and layer 4 at 118 ns against 120 ns, so the four inter-pulse intervals sum to 758 ns where the interval measured directly between the first and last pulse is 768 ns. Both deviations are in the same direction and neither exceeds one clock period. The per-layer figures come from four separate captures in which one clock period occupies a small fraction of a division, and a spread of 10 ns accumulated over four such readings is consistent with that reading uncertainty. We therefore treat the total, which is corroborated independently by the cycle counter of Section 6.8, as the authoritative figure, and the per-layer captures as confirming the model to within one cycle. Latching a counter at each `layer_done` would remove the ambiguity entirely and is the appropriate instrument for per-layer timing at this resolution.

We note that the drain term in (1) is not negligible: of the 96 cycles, 56 are spent draining the pipeline between layers rather than issuing work, so only 42 % of the core latency performs useful issue. This is a direct consequence of the ping-pong buffering, since a layer cannot begin reading a buffer until the preceding layer has finished writing it. For the narrow output layers the effect is pronounced: layer 4 issues a single token and then drains for 13 cycles. Overlapping the drain of one layer with the issue of the next would require double-buffering the activations at sub-layer granularity and is left as future work; at the present operating point the core latency is in any case a small fraction of the end-to-end path.

Arithmetic efficiency.

The consequences for throughput are worth stating explicitly, since related work is frequently compared on this basis. With 192 multipliers at 125 MHz the peak arithmetic rate is 48 GOP/s. A single inference of the reference network performs $40{\times}64 + 64{\times}32 + 32{\times}16 + 16{\times}3 = 5168$ useful multiply-accumulate operations, but the zero-padded SIMD array issues 7680 multiplications, so 67 % of the issued arithmetic is useful. Combined with the drain overhead, the achieved rate is $2 \times 5168/768\,\text{ns} \approx 13.5\,\text{GOP/s}$, or 28 % of peak. The core sustains 1.30 M inferences per second; at system level the AXI-Stream wrapper does not overlap the reception of one feature vector with the computation of the previous one, which reduces the sustained rate to approximately 0.77 M inferences per second. Both figures are far in excess of what the host-side path can supply (Section 6), which is why the design is latency-optimized rather than throughput-optimized.

The full end-to-end path was then instrumented with PS-side timestamps over 1,004,229 packets (Table 12). The result isolates where the time is spent: of the 68.6 µs mean host-to-decision latency, only 14.4 µs is spent on the board (PS + PL), and the remaining 54 µs is host-side NIC, operating-system, and wire time. On the board, the Ethernet reception (EMAC + lwIP) dominates and is bimodal (p50 2.6 µs vs p99 14.2 µs), reflecting the phase of the polling loop relative to packet arrival, whereas the DMA, PL compute, and readout stages are tightly bounded (sub-microsecond to a few microseconds, with p99 within a few percent of the median). The PL stage should be read as the interval seen by the processor rather than as the core latency: it ends when a polling read of the AXI-Lite status register first observes `done`, and one such read takes about 250 ns, so the stage is resolved in steps of that size and also contains the AXI-Stream transfer. The core latency itself is given by the cycle counter and the debug pins. This confirms the central claim: the deterministic hardware core is a negligible and jitter-free fraction of the end-to-end latency, and the remaining cost is host- and network-side software, which a hardware Ethernet MAC in the PL would remove.

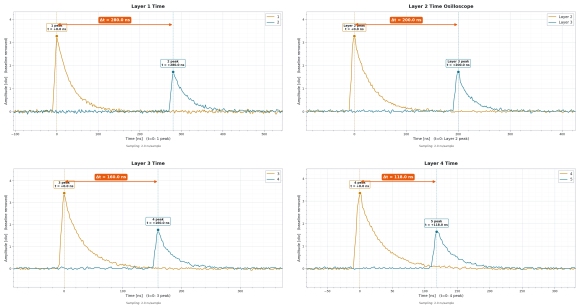


Fig. 4 Oscilloscope captures of consecutive layer-done debug pulses; each inter-pulse interval directly measures one layer's processing time. Top-left to bottom-right: layer 1 (280 ns), layer 2 (200 ns), layer 3 (160 ns), layer 4 (118 ns).

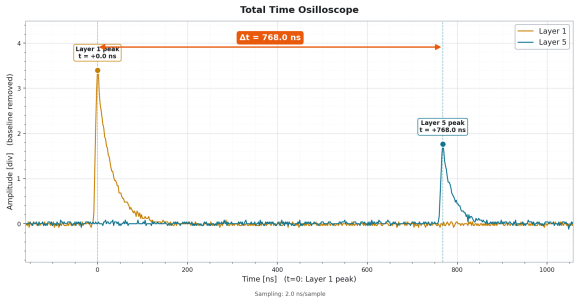


Fig. 5 Oscilloscope capture of the first and last debug pulses of an inference, measuring the total hardware-core interval ($\Delta t = 768$ ns).

6.8 Determinism of the Hardware Core

The claim that the core is deterministic has so far rested on oscilloscope captures, which show one inference at a time and cannot exclude rare outliers. To establish it over a population, a 16-bit cycle counter was added to the engine. It is cleared when `start` is accepted, increments on every cycle in which the controller is not idle, and is latched when the final layer asserts `done`. The latched value is exposed at AXI-Lite offset `0x10` and read back by the firmware after every inference, so

Table 12 Latency breakdown from Ethernet packet to trading decision, over a 1,004,229-packet run (all values in μs). Percentiles on the subtotal row are computed over the per-packet sums, not by adding the per-stage percentiles; the two differ because a packet slow in one stage is not in general slow in the others (adding the per-stage p99 values would give 21.51 rather than 21.03).

Stage	Mean	p50 / p99	Measurement Method
Ethernet packet reception (EMAC + lwIP)	7.59	2.57 / 14.18	PS timestamping (bimodal)
Packet parsing and buffer hand-off	0.70	0.69 / 0.86	PS timestamping
AXI-DMA transfer to PL	3.70	3.70 / 3.78	PS/PL counters
AXI-Stream receive + MLP core (PL)	1.46	1.46 / 1.72	`dbg` pins (core 768 ns)
Result readout	0.96	0.96 / 0.97	AXI-Lite
Board subtotal (PS + PL)	14.40	9.41 / 21.03	On-board sum
Host end-to-end	68.56	68.50 / 80.92	Host round trip

Table 13 Per-layer latency of the MLP core at 125 MHz: analytic prediction from (1) against oscilloscope measurement on the layer-done debug pins. Layers 1 to 4 are $40 \to 64$, $64 \to 32$, $32 \to 16$ and $16 \to 3$ respectively.

			Predicted		Meas.
Layer	n_ℓ	$\lceil n_\ell/3 \rceil$	cyc	ns	ns
1	64	22	36	288	280
2	32	11	25	200	200
3	16	6	20	160	160
4	3	1	15	120	118
Total			96	768	768

each logged decision carries the exact cycle count that produced it.

Over 100 000 consecutive inferences driven at line rate the counter returned 95 on every inference. Minimum, maximum and mean are all 95 and the standard deviation is zero; the distribution has a single point of support. The core therefore shows no jitter at all rather than a short tail: inference k and inference $k + 1$ occupy the same number of cycles irrespective of the input values, of the activation pattern, or of concurrent AXI-Lite and DMA traffic. This is the property the application requires, since a trading system is exposed to the worst case and not to the mean.

The counter reads 95 where (1) predicts 96 because the two count from different points. The counter increments only in non-idle cycles and is latched on the same clock edge that returns the controller to the idle state, so it excludes the cycle in which `start` is sampled; adding that cycle gives 96, in agreement with both the analytic model and the 768 ns interval measured on the debug pins (Fig. 5). The difference is a property of the instrument rather than of the datapath, and the start-to-done convention is used throughout.

6.9 Comparison Against a Serial Baseline on the Same Device

The oscilloscope measurements establish that the hardware inference completes in 768 ns, deterministically. Interpreting that number requires a baseline, and the comparisons available in the literature are awkward: the figures reported by other accelerators are measured between different points, on different devices, for different networks. We therefore use a baseline that removes all of these confounds: an earlier version of this same design, on the same board, running the same network.

That earlier implementation is preserved in the version history of the artifact and can therefore be characterized exactly rather than from recollection. It used a sequential finite-state-machine controller driving a single multiply-accumulate unit, with the same network dimensions hardwired into the controller. Its cycle count follows directly from the state machine: each neuron costs one cycle to load its bias, one cycle per input, and three further cycles to apply the activation, write the result and advance, so a layer of n_{out} neurons with n_{in} inputs occupies $n_{\text{out}}(n_{\text{in}} + 4)$ cycles. Summed over the four layers, with the layer transitions and the hardware argmax included, the design required 5639 cycles per inference.

The comparison must be stated carefully, because the two designs did not run at the same frequency: the sequential version was clocked at 100 MHz, giving 56.4 µs per inference, whereas the present design closes timing at 125 MHz. Two speed-up figures follow, and they answer different questions. Evaluated at a common 125 MHz the sequential design would require 45.1 µs, so the architectural change alone accounts for a factor of 58.7. As actually deployed, 56.4 µs then against 768 ns now, the improvement is a factor of 73.4, of which the remainder is the higher clock. We quote the former where the architecture is the subject.

The factor of 58.7 is a comparison of controller cycle counts, 5639 against 96, and it is produced by a single mechanism: replacing the single multiply-accumulate unit with 192 parallel multipliers removes the serialization of the arithmetic, so that the cost no longer scales with the 5168 multiply-accumulate operations of the network but with the number of issue tokens per layer. The earlier design also received its inputs through one AXI-Lite write per operand, which the present design replaces with a DMA-fed stream; that change shortens the path outside the core, is not contained in either cycle count, and therefore contributes nothing to the factor quoted.

This is the most defensible speed-up figure in the paper, because it is derived rather than measured under conditions we no longer control: the device, the network and the measurement points are identical, the cycle count is recoverable from the archived source, and the one confound that remains, the clock, is separated out explicitly. The cross-platform comparisons of Table 1 deserve correspondingly more caution.

6.10 Software Inference on the Same SoC

The premise of the design is that moving inference off the processor and into the fabric is worth the effort, and on this device that premise can be tested directly, since the Zynq-7000 carries a 667 MHz dual-core Cortex-A9 alongside the fabric. The same network was therefore implemented in C and run bare-metal on the PS using the same $Q8.8$ arithmetic as the hardware: the same 40-bit accumulation, the same arithmetic right shift, the same saturation, the same ReLU and the same argmax rule, with weights converted by the same routine, so that the two paths compute identical values and differ only in how the arithmetic is scheduled. Latency was measured with the Zynq Global Timer over 1000 inferences with warm caches and with no network or DMA activity inside the measured region. Each inference was additionally timed individually so that the distribution, and not only the mean, could be compared against the hardware.

The software path requires 52.4 µs per inference on average against 768 ns in the fabric, a factor of 68. Compilation settings matter here and are worth stating explicitly: the identical source built without optimization requires 302 µs, so a baseline left at the default debug settings would have overstated the speed-up by almost a factor of six. The figure reported here is compiled at `-O3` and runs on a single core with no operating system. It is scalar code: the inner product is a plain loop over 16-bit operands and NEON intrinsics were not hand-written. At 667 MHz the measured time corresponds to about 34 900 processor cycles for the 6704 multiply-accumulate operations the routine performs, or 5.2 cycles per operation, which is the cost of an unvectorized 16-bit loop on an in-order Cortex-A9. A hand-vectorized implementation would narrow the gap; we do not estimate by how much, since the rest of this paper reports measured rather than assumed quantities, and we record the absence of a NEON baseline as a limitation.

The distributions are more instructive than the means. The software path ranges from 51.2 µs at the minimum to 67.1 µs at the maximum, with a median of 52.0 µs and a 99th percentile of 52.6 µs. The spread between best and worst case is therefore 15.9 µs, which is more than twenty times the entire hardware inference, and the gap between the median and the 99th percentile alone is 570 ns, comparable to the whole of it. This is measured bare-metal, on a single core, with warm caches, with no operating system and with no network activity, which is to say under conditions more favorable than any deployment would provide; the residual variation is attributable to memory-system effects that software cannot control. The hardware path, by the counter of Section 6.8, has no spread at all. The comparison is therefore not only that the fabric is 68 times faster on average, but that its worst case equals its best case whereas the processor's worst case exceeds its best by a third of its own median.

A smaller observation points the same way. Running the identical routine with the trained weights rather than with a synthetic set changes the mean by about two percent, because the ReLU comparison is a data-dependent branch. The hardware has no counterpart to this: the pipeline evaluates the activation unconditionally, so the cycle count cannot depend on the values flowing through it.

One further comparison is worth drawing, because it sharpens what the contribution actually is. The sequential hardware design of Section 6.9 required 56.4 µs as deployed, which is slightly slower than optimized software on the same chip's processor. Moving the computation into the fabric was therefore not by itself sufficient; a serial datapath in logic bought nothing over a serial datapath in software. The gain reported here comes from the parallel issue structure, not from the change of substrate.

The comparison also has a dimension that latency alone does not express. The software path occupies the processor for the whole 52.4 µs, and during that interval it cannot service the network stack, whereas the hardware path returns the processor to its polling loop as soon as the DMA descriptor is posted. In a system whose residual latency is dominated by packet reception (Table 12), that difference compounds.

6.11 FPGA Resource Utilization

Table 14 reports the post-implementation resource usage on the XC7Z020, obtained from the routed checkpoint whose bitstream was used for the verification of Section 6.2. The core is DSP-bound (87.3 %); logic usage remains low (27.0 % of LUTs, 10.7 % of flip-flops), and 70.4 % of the block-RAM tiles are occupied. Of the totals, the inference engine itself accounts for 13 038 LUTs, 9316 flip-flops, all 192 DSP slices and 195 RAMB18 primitives; the remainder is the AXI-DMA, the interconnect and the AXI-Lite wrapper.

The block-RAM figure requires interpretation, because it does not indicate that the memories are nearly full. As described in Section 6.13, the engine instantiates one memory per (neuron-slot, input) pair. Each of these 192 memories holds 128 words of 16 bits, i.e. 2 kbit of an 18-kbit primitive. The reported occupancy therefore counts *tiles*, not stored data: roughly 88 % of the allocated block-RAM capacity is unused. Any increase in the input width or in the number of parallel neurons multiplies the tile count directly and exhausts the budget long before depth becomes a constraint.

Timing.

All constraints are met at the 125 MHz target (8.000 ns period), as summarized in Table 15: setup WNS 0.056 ns, hold WHS 0.027 ns, pulse width WPWS 3.020 ns, with zero failing endpoints out of 31,890. The corresponding maximum frequency is 125.9 MHz.

For a design whose multipliers are fully pipelined through the DSP input, intermediate and output registers, this figure is lower than one would expect, and the critical path explains why. The worst path does not lie in the arithmetic at all: it runs from a payload register inside the AXI SmartConnect to the data input of one of the 192 weight memories, with a data-path delay of 7.532 ns of which 7.014 ns (93.1 %) is routing and only 0.518 ns is logic, with no logic level in between. The limiting factor is thus the physical distribution of the AXI-Lite write path across a large number of small, spatially dispersed memories, which is the same organizational choice identified above as the scaling limit. The worst path that starts and ends inside the engine has a slack of 0.113 ns and has the same character: it runs from a write-back register to one of the 64 entries of the activation buffer, 91.6 % of its delay being routing across a single LUT. Both paths are fan-out paths that move a value to many dispersed destinations; neither passes through the multiply-reduce datapath. Consolidating inputs into wider memory words would shorten the first of them and raise the achievable frequency, in addition to relaxing the capacity bound; we did not pursue it here because 125 MHz is already supplied directly by the PS clock generator and the core latency is not the dominant term in the end-to-end path (Section 6).

Power.

Vivado's vectorless estimate for the implemented design is 2.704 W of total on-chip power (2.509 W dynamic, 0.196 W static) at a junction temperature of 56.2 °C. We emphasize that this is a tool estimate obtained by activity propagation rather than a measurement. The breakdown in Table 16 shows that the processing system accounts for 1.535 W, i.e. more than half of the total, while the entire programmable-logic contribution (clocks, logic, signals, block RAM, DSP and I/O) amounts to 0.974 W. Comparisons against accelerator cards should use the latter figure, since those designs do not include an application processor in their reported power.

6.12 Deployment Cost

The platform used here is a Digilent ZedBoard carrying an XC7Z020, a part in the low hundreds of US dollars, against the several thousand typical of a data-center accelerator card. We deliberately avoid a tabulated cost comparison against CPU, GPU and

Table 14 Post-implementation resource utilization on the Xilinx XC7Z020-1CLG484

Resource	Used	Available	Utilization
Slice LUTs	14,366	53,200	27.00%
as logic	14,282	53,200	26.85%
as memory	84	17,400	0.48%
Slice registers	11,396	106,400	10.71%
Block RAM tiles	98.5	140	70.36%
DSP slices	192	220	87.27%
Bonded IOB	14	200	7.00%
BUFGCTRL	1	32	3.13%

Table 15 Post-implementation timing closure at 125 MHz (8.000 ns period, speed grade −1)

Check	Worst slack	Failing endpoints
Setup (WNS)	0.056 ns	0 / 31,890
Hold (WHS)	0.027 ns	0 / 31,890
Pulse width (WPWS)	3.020 ns	0 / 12,115

Table 16 On-chip power, Vivado vectorless estimate for the implemented design

Component	Power
PS7 (processing system)	1.535 W
Signals	0.272 W
Block RAM	0.249 W
DSP	0.236 W
Slice logic	0.145 W
Clocks	0.067 W
I/O	0.005 W
Dynamic total	2.509 W
Device static	0.196 W
Total on-chip	2.704 W
of which PL only	0.974 W

high-end FPGA alternatives, because we have not measured this network on those platforms and a table populated with qualitative labels would add nothing. What the measurements do support is narrower and can be stated directly: a deterministic, sub-microsecond inference path for a network of this size does not require a data-center part. As a point of reference, SparkNoC [19] reaches 337.2 GOP/s at about 7.6 W of compute power on an Arria 10 GX1150 card, an order of magnitude more throughput than the present design delivers, on hardware of a different class and for a throughput-oriented workload rather than single-inference latency.

6.13 Run-Time Reconfigurability and Scalability

The framework property claimed in Section 3 is that a single bitstream hosts different networks. This is not a matter of reloading weights alone: the per-layer output widths are held in AXI-Lite configuration registers that the issue FSM reads to derive its iteration counts, so the *shape* of the network is also a run-time parameter.

The verification campaign of Section 6.2 demonstrates this directly. The four networks of Table 8 have different hidden-layer widths, with the first hidden layer ranging from 10 to 64 neurons, yet all four were loaded into the same programmed device over AXI-Lite and each was verified bit-exactly. No re-synthesis, re-implementation or bitstream reload took place between them; loading a network requires $\sum_\ell \lceil n_\ell/3 \rceil \times 3 \times 64$ weight writes plus the corresponding biases, which for the reference topology is approximately 7800 AXI-Lite transactions.

The latency follows the shape rather than the bounds of the datapath. Because all 64 inputs of a neuron slot are multiplied in parallel, the input width of a layer does not enter the cycle count; only the number of output neurons does, through the $\lceil n_\ell/3 \rceil$ issue term of (1). The four networks therefore complete in 69, 76, 95 and 111 counted cycles (Table 8), i.e. 560 ns to 896 ns start to done, and each count was constant over its 100 000 inferences. The price of hosting several shapes in one datapath is visible in the same expression: the drain term of 4×14 cycles is independent of the shape, so it is more than half of the latency of the smaller networks, and a network much smaller than the bounds gains less than its arithmetic would suggest.

The limits of this reconfigurability should be stated precisely. The datapath is elaborated for at most four layers, at most 64 neurons per layer and at most 64 inputs per layer; these bounds are compile-time constants of the SIMD array and of the weight-memory instantiation. Any network within those bounds is hosted without re-synthesis; a network exceeding them is not.

Scaling beyond these bounds is limited by memory organization rather than by arithmetic. The design instantiates one memory per (neuron-slot, input) pair, i.e. PARALLEL_NEURONS × MAX_INPUTS = 192 weight memories, each of which stores only 128 words and therefore occupies a small fraction of the 18-kbit

primitive it is mapped onto. The reported 70.4 % block-RAM occupancy thus reflects the *number* of memories rather than their depth: widening the input vector or increasing the number of parallel neurons increases the block count proportionally and exhausts the budget well before the depth of any individual memory becomes a constraint. Packing several inputs per memory word is the natural way to relax this limit and is left as future work.

6.14 Comparison with a Compiler-Generated Implementation

The figures of Table 1 are measured on different devices, for different networks and between different points. To place a measured point at the compiler-based end of the design space, the four networks of Table 8 were also generated with hls4ml 1.3.0 [14] for the same part (XC7Z020-1CLG484) at the same 8 ns target, using Vitis HLS and Vivado 2025.2.

Same arithmetic.

The hls4ml models were configured to reproduce the arithmetic of the engine rather than hls4ml's defaults: weights, biases and activations in `ap_fixed<16,8>`, accumulation in `ap_fixed<40,24>` so that no product or partial sum is rounded, and a result type with truncation and saturation, which matches the arithmetic shift and 16-bit clamp of the engine. Weights were pre-quantized with the same truncating conversion as the on-board loader. Before synthesis, the C emulation of each generated model was compared with the golden model of Section 6.2 on 20 000 inputs of the verification stimulus: all four networks agreed on every raw output and every decision. The comparison is therefore between two implementations of the same function, not between a 16-bit design and a narrower one.

Design points.

hls4ml trades area for latency through the reuse factor (RF), the number of multiplications each multiplier performs per inference. Each network was generated with the `Resource` strategy and a parallel interface at RF $\in \{8, 16, 32, 64, 128\}$, giving 20 designs. Each was synthesized, and every design that fits the device was then placed and routed out of context with an 8 ns constraint on its clock, which yields a post-route worst negative slack and hence an achievable frequency. The synthesis estimate of HLS was not used for this purpose: on the small networks the points it estimated at 5.75 ns routed at about 7.8 ns, and on the large ones the same estimate routed at up to 11 ns, so an estimate-based comparison would have credited hls4ml with timing closure it does not achieve. Table 17 reports the result.

Result.

Only the four designs marked in Table 17 close timing at 125 MHz, three of hls4ml and the engine. On the two smaller networks, hls4ml at RF 8 reaches 456 to 457 ns against 560 and 616 ns for the engine, i.e. it is 1.2 to 1.35 times faster, at a cost of 84 and 201 DSP slices for a single network; at RF 16 on 40-32-16-8-3 the two are within 6 % of each other. On the two larger networks no hls4ml point meets 125 MHz. The designs with enough parallelism to compete either exceed the LUT capacity of the device (RF 8, and RF 16 for 40-64-64-32-3) or fit but route at 91 to 125 MHz. Even when each design is credited with its own achievable frequency, the fastest hls4ml implementation of 40-64-32-16-3 takes 1026 ns against 768 ns, and the fastest of 40-64-64-32-3 takes 1758 ns against 896 ns.

Three qualifications apply, and two of them favour hls4ml. First, the hls4ml designs were routed out of context, without the processing system, the DMA and the interconnect that share the device with the engine, and without clock insertion delay; the engine's figures come from the complete system. Second, hls4ml pipelines successive inferences, accepting a new input every II cycles, so its throughput exceeds that of the engine by a large factor. The engine cannot overlap inferences in this way, because the same 192 DSP slices evaluate every layer in turn: a second input can enter only once the first has left the last layer. hls4ml obtains the overlap by giving each layer its own multipliers, which is what its DSP counts in Table 17 pay for. For the single-inference latency that this application is exposed to, the overlap does not help, and the throughput gap does not reach the system: even the engine's 1.30 M inferences per second is roughly two orders of magnitude above the 11 000 to 12 400 requests per second that the Linux host sustained in the runs of Table 11, with one request in flight at a time. Third, each hls4ml row is a separate bitstream for one network with its weights compiled into the logic, so moving between the networks of Table 17, or retraining any of them, requires a new synthesis and place-and-route run, which took from a few minutes to over half an hour per point here. The engine runs all four from one bitstream, at a fixed 192 DSP slices, and changes network with AXI-Lite writes.

Table 17 hls4ml implementations of the four networks on the XC7Z020 at an 8 ns target, against this engine. Latency is the HLS worst-case latency in cycles; f_{max} is obtained from the post-route worst negative slack, and * marks designs that close timing at 125 MHz; "Time" is latency at f_{max}. "Does not fit" marks designs whose synthesis exceeds the 53 200 LUTs of the device. The engine is a single bitstream that hosts all four networks; its cycle count is start to done and its LUT figure is for the engine alone.

Network	RF	Latency (cyc)	II (cyc)	DSP	LUT	f_{max} (MHz)	Time (ns)
40-10-16-8-3	8	59	10	84	5,624	129.0*	457
	16	87	20	39	7,045	118.5	734
	32	158	40	19	7,358	117.6	1,344
	64	270	80	10	8,044	97.6	2,766
	128	414	160	8	8,828	79.6	5,201
40-32-16-8-3	8	58	8	201	10,105	127.3*	456
	16	84	20	108	9,000	128.5*	654
	32	150	40	53	9,322	113.1	1,326
	64	254	80	27	13,941	101.9	2,493
	128	462	160	14	12,096	98.8	4,676
40-64-32-16-3	8	60	8	220	119,758	does not fit	
	16	93	20	220	31,940	90.6	1,026
	32	143	40	148	16,721	124.8	1,146
	64	278	80	73	23,943	105.7	2,630
	128	486	160	37	24,066	103.6	4,691
40-64-64-32-3	8	61	8	219	228,808	does not fit	
	16	94	20	220	91,808	does not fit	
	32	160	40	220	31,601	91.0	1,758
	64	262	80	131	28,935	103.5	2,531
	128	534	160	65	34,651	110.3	4,841
This work, all four	–	70 / 77 / 96 / 112	≥ lat.	192	13,038	125.9*	560 / 616 / 768 / 896

7 Discussion

The results support four observations. First, embedding the network in the PL makes the core latency a property of the pipeline rather than of the operating system: the measured core latency is 768 ns and deterministic, whereas the end-to-end host path around it varies from 66 to 460 µs in median across operating systems and shows heavy tails. Second, $Q8.8$ quantization is essentially accuracy-neutral here (at most 0.22 percentage points across five models), so the speed-up does not cost predictive quality. Third, the design fits a low-cost Zynq-7000 at 2.70 W with positive timing slack, so the approach does not depend on premium accelerator hardware. Fourth, run-time shape does not cost latency against a compiler that specializes the hardware to each network: on the device used here, hls4ml is faster only on the two smaller networks and only at its highest-parallelism points, while on the two larger networks the engine is 1.3 to 2.0 times faster than the fastest hls4ml design that fits (Section 6.14). What the specialization does buy is throughput, through inference-level pipelining that the engine does not implement. The instrumented per-stage breakdown (Table 12) further shows that only 14.4 µs of the 68.6 µs end-to-end latency is spent on the board, so the main remaining limitation is the host- and network-side software path; moving the Ethernet stack into the PL is the clearest next step.

8 Conclusion

We presented an Ethernet-fed MLP inference framework embedded in the programmable logic of a low-cost Zynq-7000 SoC. Feature vectors are streamed to the PL over AXI-DMA, and a 192-multiplier, 13-stage pipelined core performs the $Q8.8$ inference in 768 ns with deterministic timing, using 27.0 % of the LUTs and 87.3 % of the DSP slices of an XC7Z020 at an estimated 2.70 W of total on-chip power, of which 0.97 W is attributable to the programmable logic. The same network executed bare-metal on the SoC's own Cortex-A9 requires 52.4 µs, so moving the computation into the fabric is worth a factor of 68 against an independent baseline on identical arithmetic

(Section 6.10). Decomposed against an earlier sequential implementation of the same network on the same device, the architecture accounts for a factor of 58.7 at a common clock, and 73.4 as the two were actually deployed (Section 6.9).

The property that distinguishes the design from compiler-style FPGA inference frameworks is that the network shape is a run-time rather than an elaboration-time parameter: four networks with different hidden-layer widths were deployed into one unchanged bitstream over AXI-Lite, with no synthesis step between them, and the cycle counter confirms that each ran at its own shape. Against hls4ml implementations of the same four networks on the same device, clock and arithmetic, this flexibility does not cost single-inference latency on the larger networks, where no hls4ml design closes timing at 125 MHz. Correctness of the deployed arithmetic is not assumed but established: the core was verified bit-exactly against an independent reference model over 400,000 inferences, a process that first exposed and then eliminated three pipeline-alignment defects that neither synthesis nor end-to-end testing had revealed.

Two directions follow from the measured breakdown. First, only 14.4 µs of the 68.6 µs end-to-end latency is spent on the board, so integrating an Ethernet MAC directly in the programmable logic, and thereby removing the host network stack from the path, is the change with the largest remaining effect. Second, the present bounds on network shape (at most four layers of at most 64 neurons and 64 inputs) follow from the one-memory-per-input organization of the weight storage rather than from the arithmetic; packing several inputs per memory word would relax them without additional DSP cost.

Statements and Declarations

Funding. This research did not receive any specific grant from funding agencies in the public, commercial, or not-for-profit sectors.

Competing interests. The authors declare that they have no known competing financial interests or personal relationships that could have appeared to influence the work reported in this paper.

Ethics approval and consent to participate. Not applicable.

Consent for publication. Not applicable.

Data availability. The market data used in this study are publicly available Binance BTCUSDT aggregated trade records. The measurement logs and verification stimuli are part of the artifact described under Code availability.

Code availability. The artifact supporting this study is publicly available at https://github.com/ozkanerbatuhan/hft (hardware) and https://github.com/ozkanerbatuhan/HFT_vitis (firmware and host tooling). It contains the RTL sources, the self-checking testbench, the cycle-accurate reference model, the host and firmware tooling, the stimulus generator, the post-implementation reports, and the measurement logs behind the latency and verification results. Instructions for rebuilding the bitstream and for reproducing each claim are given in the accompanying documentation.

Author contributions. Ahmet Batuhan Özkaner: conceptualization, methodology, hardware and software implementation, verification, formal analysis, investigation, data curation, visualization, and writing of the original draft. Mustafa Alper Ilık: software, investigation, and review and editing of the manuscript. Baykal Sarıoğlu: supervision and project administration. All authors read and approved the final manuscript.

Use of AI tools. Large language model based assistants were used as an aid in the literature search, in the development of the measurement and verification software, and for correcting typographical and language errors in the manuscript. The authors reviewed all such output and take full responsibility for the content of this article.

References

[1] Budish, E., Cramton, P., Shim, J.: The high-frequency trading arms race: Frequent batch auctions as a market design response. The Quarterly Journal of Economics **130**(4), 1547–1621 (2015) https://doi.org/10.1093/qje/qjv027

[2] Aquilina, M., Budish, E., O'Neill, P.: Quantifying the high-frequency trading "arms race". The Quarterly Journal of Economics **137**(1), 493–564 (2022) https://doi.org/10.1093/qje/qjab032

[3] Leber, C., Geib, B., Litz, H.: High frequency trading acceleration using FPGAs. In: Proceedings of the 21st International Conference on

Field Programmable Logic and Applications (FPL), pp. 317–322 (2011). https://doi.org/10.1109/FPL.2011.64

[4] Lockwood, J.W., Monga, M.: Implementing ultra low latency data center services with programmable logic. In: Proceedings of the IEEE 23rd Annual Symposium on High-Performance Interconnects (HOTI), pp. 68–77. IEEE, ??? (2015). https://doi.org/10.1109/HOTI.2015.20

[5] Puranik, S., Barve, M., Rodi, S., Patrikar, R.: Acceleration of trading system back end with FPGAs using high-level synthesis flow. Electronics **12**(3), 520 (2023) https://doi.org/10.3390/electronics12030520

[6] Brown, N., Klaisoongnoen, M., Brown, O.T.: Optimisation of an FPGA credit default swap engine by embracing dataflow techniques. In: IEEE International Conference on Cluster Computing (CLUSTER), FPGA for HPC Workshop (2021). https://doi.org/10.48550/arXiv.2108.03982

[7] Yoo, S., Kim, H., Kim, J., Park, S., Kim, J.-Y., Oh, J.: LightTrader: A standalone high-frequency trading system with deep learning inference accelerators and proactive scheduler. In: Proceedings of the IEEE International Symposium on High-Performance Computer Architecture (HPCA), pp. 1017–1030 (2023). https://doi.org/10.1109/HPCA56546.2023.10070930

[8] Zhang, C., Li, P., Sun, G., Guan, Y., Xiao, B., Cong, J.: Optimizing FPGA-based accelerator design for deep convolutional neural networks. In: Proceedings of the ACM/SIGDA International Symposium on Field-Programmable Gate Arrays (FPGA), pp. 161–170 (2015). https://doi.org/10.1145/2684746.2689060

[9] Qiu, J., Wang, J., Yao, S., Guo, K., Li, B., Zhou, E., Yu, J., Tang, T., Xu, N., Song, S., Wang, Y., Yang, H.: Going deeper with embedded FPGA platform for convolutional neural network. In: Proceedings of the ACM/SIGDA International Symposium on Field-Programmable Gate Arrays (FPGA), pp. 26–35 (2016). https://doi.org/10.1145/2847263.2847265

[10] Guo, K., Sui, L., Qiu, J., Yu, J., Wang, J., Yao, S., Han, S., Wang, Y., Yang, H.: Angel-Eye: A complete design flow for mapping CNN onto embedded FPGA. IEEE Transactions on Computer-Aided Design of Integrated Circuits and Systems **37**(1), 35–47 (2018) https://doi.org/10.1109/TCAD.2017.2705069

[11] Umuroglu, Y., Fraser, N.J., Gambardella, G., Blott, M., Leong, P., Jahre, M., Vissers, K.: FINN: A framework for fast, scalable binarized neural network inference. In: Proceedings of the ACM/SIGDA International Symposium on Field-Programmable Gate Arrays (FPGA), pp. 65–74 (2017). https://doi.org/10.1145/3020078.3021744

[12] Blott, M., Preußer, T.B., Fraser, N.J., Gambardella, G., O'Brien, K., Umuroglu, Y., Leeser, M., Vissers, K.: FINN-R: An end-to-end deep-learning framework for fast exploration of quantized neural networks. ACM Transactions on Reconfigurable Technology and Systems **11**(3), 1–23 (2018) https://doi.org/10.1145/3242897

[13] Duarte, J., Han, S., Harris, P., Jindariani, S., Kreinar, E., Kreis, B., Ngadiuba, J., Pierini, M., Rivera, R., Tran, N., Wu, Z.: Fast inference of deep neural networks in FPGAs for particle physics. Journal of Instrumentation **13**(07), 07027 (2018) https://doi.org/10.1088/1748-0221/13/07/P07027

[14] Schulte, J.-F., Ramhorst, B., Sun, C., Mitrevski, J., Ghielmetti, N., Lupi, E., *et al.*: hls4ml: A flexible, open source platform for deep learning acceleration on reconfigurable hardware. ACM Transactions on Reconfigurable Technology and Systems **19**(2), 1–35 (2026) https://doi.org/10.1145/3801979

[15] Sun, C., Que, Z., Loncar, V., Luk, W., Spiropulu, M.: da4ml: Distributed arithmetic for real-time neural networks on FPGAs. ACM Transactions on Reconfigurable Technology and Systems **19**(1), 1–27 (2026) https://doi.org/10.1145/3777387

[16] Broglie, G., Morge-Rollet, L., Le Jeune, D., Le Roy, F., Roland, C., Canaff, C., Diguet, J.-P.: New LTE-A low-latency detectors for cognitive radio on low-cost FPGA SoC. Journal of Signal Processing Systems **96**(8-9), 479–493 (2024)

https://doi.org/10.1007/s11265-024-01932-9

[17] Gdaim, S., Bensalem, H., Gagnon, G.: Efficient FPGA acceleration of deep neural networks for image processing using an OpenCL-based high-level synthesis framework. Journal of Signal Processing Systems **98**(2), 51 (2026) https://doi.org/10.1007/s11265-026-02014-8

[18] Rangsikunpum, A., Amiri, S., Ost, L.: Low-cost FPGA-based hybrid defence against sticker attacks for traffic-sign classification. Journal of Signal Processing Systems **98**(1), 20 (2026) https://doi.org/10.1007/s11265-026-01994-x

[19] Xia, M., Huang, Z., Tian, L., Wang, H., Chang, V., Zhu, Y., Feng, S.: SparkNoC: An energy-efficiency FPGA-based accelerator using optimized lightweight CNN for edge computing. Journal of Systems Architecture **115**, 101991 (2021) https://doi.org/10.1016/j.sysarc.2021.101991

[20] Tsoutsouras, V., Koliogeorgi, K., Xydis, S., Soudris, D.: An exploration framework for efficient high-level synthesis of support vector machines: Case study on ECG arrhythmia detection for Xilinx Zynq SoC. Journal of Signal Processing Systems **88**(2), 127–147 (2017) https://doi.org/10.1007/s11265-017-1230-1

[21] Wang, T., Wang, C., Zhou, X., Chen, H.: A survey of FPGA based deep learning accelerators: Challenges and opportunities. arXiv preprint arXiv:1901.04988 (2019) https://doi.org/10.48550/arXiv.1901.04988

[22] Rios-Navarro, A., Tapiador-Morales, R., Jimenez-Fernandez, A., Dominguez-Morales, M., Amaya, C., Linares-Barranco, A.: Performance evaluation over HW/SW co-design SoC memory transfers for a CNN accelerator. In: Proceedings of the IEEE International Conference on Electronics, Circuits and Systems (ICECS) (2018). https://doi.org/10.48550/arXiv.1806.01106

[23] Jacob, B., Kligys, S., Chen, B., Zhu, M., Tang, M., Howard, A., Adam, H., Kalenichenko, D.: Quantization and training of neural networks for efficient integer-arithmetic-only inference. In: Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR), pp. 2704–2713 (2018). https://doi.org/10.1109/CVPR.2018.00286

[24] Hubara, I., Courbariaux, M., Soudry, D., El-Yaniv, R., Bengio, Y.: Quantized neural networks: Training neural networks with low precision weights and activations. Journal of Machine Learning Research **18**(1), 6869–6898 (2017)

[25] Gholami, A., Kim, S., Dong, Z., Yao, Z., Mahoney, M.W., Keutzer, K.: A survey of quantization methods for efficient neural network inference. arXiv preprint arXiv:2103.13630 (2021) https://doi.org/10.48550/arXiv.2103.13630

[26] Al-Rikabi, H., Renczes, B.: Floating-point quantization analysis of multi-layer perceptron artificial neural networks. Journal of Signal Processing Systems **96**(4-5), 301–312 (2024) https://doi.org/10.1007/s11265-024-01911-0

[27] Moreau, T., Chen, T., Vega, L., Roesch, J., Yan, E.Q., Zheng, L., Fromm, J., Jiang, Z., Ceze, L., Guestrin, C., Krishnamurthy, A.: A hardware–software blueprint for flexible deep learning specialization. IEEE Micro **39**(5), 8–16 (2019) https://doi.org/10.1109/MM.2019.2928962

[28] Yu, Y., Wu, C., Zhao, T., Wang, K., He, L.: OPU: An FPGA-based overlay processor for convolutional neural networks. IEEE Transactions on Very Large Scale Integration (VLSI) Systems **28**(1), 35–47 (2020) https://doi.org/10.1109/TVLSI.2019.2939726

[29] Yu, Y., Zhao, T., Wang, K., He, L.: Light-OPU: An FPGA-based overlay processor for lightweight convolutional neural networks. In: Proceedings of the ACM/SIGDA International Symposium on Field-Programmable Gate Arrays (FPGA), pp. 122–132 (2020). https://doi.org/10.1145/3373087.3375311

[30] AMD: Vitis AI Development Environment. Product documentation (2024). https://www.amd.com/en/developer/resources/vitis-ai.html

[31] Khaki, A.M.Z., Choi, A.: Real-time and resource-efficient embedded computer vision via optimizing lightweight CNNs for FPGA acceleration. Journal of Signal Processing Systems **97**(2–4), 185–195 (2025) https://doi.org/10.1007/s11265-025-01966-7

[32] Kabir, E., Bakos, J.D., Andrews, D., Huang, M.: A runtime-adaptive transformer neural network accelerator on FPGAs. Microprocessors and Microsystems **120**, 105223 (2026) https://doi.org/10.1016/j.micpro.2025.105223

[33] Ratto, F., Máinez, A.P., Sau, C., Meloni, P., Deriu, G., Delucchi, S., Massa, M., Raffo, L., Palumbo, F.: An automated design flow for adaptive neural network hardware accelerators. Journal of Signal Processing Systems **95**(9), 1091–1113 (2023) https://doi.org/10.1007/s11265-023-01855-x

[34] Ezzaoui Rahali, H., Dave, A., Ruckman, L., Rahimifar, M.M., Therrien, A.C., Russel, J.J., Herbst, R.T.: Neural Network Acceleration on MPSoC Board: Integrating SLAC's SNL, Rogue Software and Auto-SNL. arXiv preprint arXiv:2508.21739 (2025). https://doi.org/10.48550/arXiv.2508.21739

[35] Boudjadar, J., Islam, S.U., Buyya, R.: Dynamic FPGA reconfiguration for scalable embedded artificial intelligence (AI): A co-design methodology for convolutional neural networks (CNN) acceleration. Future Generation Computer Systems **169**, 107777 (2025) https://doi.org/10.1016/j.future.2025.107777

[36] Umuroglu, Y., Akhauri, Y., Fraser, N.J., Blott, M.: LogicNets: Co-designed neural networks and circuits for extreme-throughput applications. In: Proceedings of the 30th International Conference on Field Programmable Logic and Applications (FPL), pp. 291–297 (2020). https://doi.org/10.1109/FPL50879.2020.00055

[37] Andronic, M., Constantinides, G.A.: PolyLUT: Learning piecewise polynomials for ultra-low latency FPGA LUT-based inference. In: Proceedings of the International Conference on Field Programmable Technology (ICFPT), pp. 60–68 (2023). https://doi.org/10.1109/ICFPT59805.2023.00012

[38] Andronic, M., Constantinides, G.A.: NeuraLUT: Hiding neural network density in boolean synthesizable functions. In: Proceedings of the 34th International Conference on Field-Programmable Logic and Applications (FPL), pp. 140–148 (2024). https://doi.org/10.1109/FPL64840.2024.00028

[39] Lou, B., Rademacher, R., Boland, D., Leong, P.H.W.: PolyLUT-Add: FPGA-based LUT inference with wide inputs. In: Proceedings of the 34th International Conference on Field-Programmable Logic and Applications (FPL), pp. 149–155 (2024). https://doi.org/10.1109/FPL64840.2024.00029

[40] Andronic, M., Constantinides, G.A.: NeuraLUT-Assemble: Hardware-aware assembling of sub-neural networks for efficient LUT inference. In: Proceedings of the IEEE 33rd Annual International Symposium on Field-Programmable Custom Computing Machines (FCCM), pp. 208–216 (2025). https://doi.org/10.1109/FCCM62733.2025.00077

[41] Khataei, A., Bazargan, K.: TreeLUT: An efficient alternative to deep neural networks for inference acceleration using gradient boosted decision trees. In: Proceedings of the ACM/SIGDA International Symposium on Field Programmable Gate Arrays (FPGA), pp. 14–24 (2025). https://doi.org/10.1145/3706628.3708877

[42] Zhang, Z., Zohren, S., Roberts, S.: DeepLOB: Deep convolutional neural networks for limit order books. IEEE Transactions on Signal Processing **67**(11), 3001–3012 (2019) https://doi.org/10.1109/TSP.2019.2907260

[43] Ntakaris, A., Magris, M., Kanniainen, J., Gabbouj, M., Iosifidis, A.: Benchmark dataset for mid-price forecasting of limit order book data with machine learning methods. Journal of Forecasting **37**(8), 852–866 (2018) https://doi.org/10.1002/for.2543

[44] Gutermann, A., Serdyuk, A., Paraskevas, F., Toto Kiesa, H., Lesniak, F., Schwarz, J., Hartmann, M., Harbaum, T., Becker, J.: Joint exploration of neural networks and systolic hardware for improved AI accelerator performance. Journal of Signal Processing Systems **98**(2), 42 (2026)

https://doi.org/10.1007/s11265-026-02012-w

Author Biographies

Ahmet Batuhan Özkaner. received the B.Sc. degree in Electrical and Electronics Engineering from İstanbul Bilgi University, İstanbul, Türkiye. He is currently an electrical design engineer at Rönesans Holding. His research interests include FPGA-based digital design, hardware acceleration of neural network inference, and embedded systems.

Mustafa Alper Ilık. is with the Department of Electrical and Electronics Engineering, İstanbul Bilgi University, İstanbul, Türkiye. His research interests include embedded systems and FPGA-based hardware design.

Baykal Sarıoğlu. is a faculty member in the Department of Electrical and Electronics Engineering, İstanbul Bilgi University, İstanbul, Türkiye. His research interests include electronic circuits and systems.